

tomcarron__exercise9

June 10, 2022

1 Exercise Set 9

Due: **9:30 13 June 2022**

Discussion: **13:00 17 June 2022**

Online submission at via [ILIAS](#) in the directory Exercises / Übungen -> Submission of Exercises / Rückgabe des Übungsblätter

Data analysis and code development is typically done in larger collaborations. The most common way to do this is with a cloud service such as [GitHub](#). `git` is a common method used for version control. You can find a descriptions of the various capabilities for example at [Atlassian](#).

It is recommended that you register on GitHub to get used to it. For the remainder of the course, you may keep your solutions on GitHub and simply send a link to the necessary folder as the ILIAS submission (for example in an ascii file). If you choose to do so, the requirements for a submission remain the same (submit a self-contained pdf) and additional requirements that your repository is called `DataAnalysis` and the exercise solutions are placed in directories `exercise_x/` (so for exercise 9 it would be `exercise_9/`).

```
[ ]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import os
import seaborn as sb
from sklearn import linear_model
from scipy.optimize import curve_fit
from matplotlib import collections as matcoll
import scipy.stats as st
import matplotlib
matplotlib.rcParams['figure.dpi']=300
```

2 1. Regression - Fitting a line [data exercise] [50 Points]

In this problem we will fit a linear function $f(x) = y_0 + \beta_0 x$ to the data from datafile: `line_data.dat`.


a) Perform the linear regression to fit a linear function $f(x) = y_0 + \beta_0 x$ and plot $f(x)$ together with the data points. **25 Points**

```
[ ]: # load the data
df = pd.read_csv("line_data.dat", sep="\s+", names=["x", "y"])
# as arrays
x, y = df["x"].to_numpy(), df["y"].to_numpy()
# Function for fitting
def lin_func(x, y0, beta):
    f = y0 + beta * x
    return f

# linear fit
popt, pcov = curve_fit(lin_func, x, y)
x2 = np.linspace(np.min(x), np.max(x), 1000)
# difference between data and the fit
y_predicted = lin_func(x, popt[0], popt[1])
diff = y - y_predicted

# lines between prediction and data
lines = []
for i in range(len(x)):
    pair = [(x[i], y[i]), (x[i], y_predicted[i])]
    lines.append(pair)

linecoll = matcoll.LineCollection(lines, colors="k")
# plot the data
fig, ax = plt.subplots()
ax.scatter(x, y, label="data", s=9)
# ax.
    → scatter(x, lin_func(x, popt[0], popt[1]), s=10, c='green', marker='o', label='predicted')
ax.plot(x2, lin_func(x2, popt[0], popt[1]), ls="--", c="red", label="fit")
ax.add_collection(linecoll)
ax.set_xlabel("x")
ax.set_ylabel("y")
ax.legend()
plt.savefig("plots/linreg.png", dpi=400, bbox_inches="tight")
```




b) Compute the residuals and plot them. What is R^2 ? **25 Points**

```
[ ]: # Compute the residuals and plot them
# lines between prediction and data
lines = []
for i in range(len(x)):
    pair = [(x[i], 0), (x[i], y[i] - y_predicted[i])]
    lines.append(pair)

linecoll = matplotlib.LineCollection(lines, colors="k")

fig, ax = plt.subplots()
ax.scatter(x, diff, label="data", s=9)
ax.plot(x2, np.zeros_like(x2), "--", c="r")
ax.add_collection(linecoll)
ax.set_ylabel("Residuals")
ax.set_xlabel("x")
ax.legend()
plt.savefig("plots/residuals.png", dpi=400, bbox_inches="tight")
```



```
[ ]: # R^2
# In the case of only one dependent variable, R^2 is the square of the pearson
↪ corr coeff
R2 = (st.pearsonr(x, y)[0]) ** 2
print("R2 = ", R2)
```

R2 = 0.6124158502151285

3 2. Model fit quality assessment [50 Points]

In this problem you will assess the quality of a fit by inspecting the residuals of the fit. For each of the following plots state whether the residuals indicate a reasonable model fit and briefly explain your conclusion and if applicable how possibly to improve the fit. **50 Points**

4 Plot a)

The residuals indicate a good model fit because the residuals are centred around $y=0$ and are normally distributed.

5 Plot b)

The residuals of plot b also indicate a good model fit because the residuals are centred around $y=0$ and are normally distributed.

6 Plot c)

The residuals show a trend with a negative autocorrelation. This suggests the model is not suitable for the data.

7 Plot d)

The residuals are normally distributed about $y=0$. The model is a good fit.

8 Plot e)

Trend in the residuals, the model is not suitable for the data. A polynomial model would fit the data better. there is also some negative autocorrelation.